

Lecture 12: OpenMP

Producers/Consumers – Queues

- An abstraction of a line of customers waiting to pay for their groceries in a supermarket
- A natural data structure to use in many multithreaded applications
- For several "producer" and "consumer" threads
 - Producer threads might "produce" requests for data
 - Consumer threads might "consume" the request by finding or generating the requested data

Message-Passing

- Each thread could have a shared message queue, and when one thread wants to "send a message" to another thread, it could **enqueue** the message in the destination thread's queue
- A thread could receive a message by **dequeuing** the message at the head of its message queue

Message-Passing: Pseudo Code for Each Thread

```
for (sent_msgs = 0; sent_msgs < send_max ; sent_msgs++  
{  
    Send_msg();  
    Try_receive ();  
}
```

```
while (! Done ())  
    Try_receive ();
```

Pseudo Code for Send_msg

```
mesg = random( );  
dest = random( ) % thread_count ;  
# pragma omp critical  
Enqueue(queue, dest, my_rank, mesg);
```

Receiving Messages

- Only the destination queue-owner thread can **Dequeue** the enqueued message
- If the queue contains only a single message then **Dequeue** must be performed in a **critical section** so that no other thread can call Enqueue
- If the queue contains at least two messages then Dequeue can be performed without the use of a critical section

Receiving Messages

- A correct implementation of the previous synchronization scheme can be achieved by introducing shared variables **enqueued** and **dequeued** that counts the number of messages sent and received, respectively

Pseudo Code for Try_receive

- By introducing the derived variable
queue_size = enqueued – dequeued
the **Try_receive** function can be implemented as:

```
queue_size = enqueued – dequeued ;  
if (queue_size == 0) return ;  
else if (queue_size == 1)  
# pragma omp critical  
    Dequeue(queue, &src, &mesg) ;  
else  
    Dequeue(queue, &src, &mesg) ;  
Print_message(src, mesg) ;
```


Pseudo Code for Try_receive

- Synchronization issues: recall **enqueued** and **dequeued** are shared variables

- Suppose `queue_size = 0` or `1` when the actual values should be `1` or `2`
- If `queue_size = 0`, `Try_receive` returns but tries again later
- If `queue_size = 1`, `Dequeue` is called under a critical section but `Try_receive` is called again if the actual `queue_size` should have been `2`

```
queue_size = enqueued - dequeued ;  
if (queue_size == 0) return ;  
else if (queue_size == 1)  
# pragma omp critical  
    Dequeue(queue, &src, &mesg) ;  
else  
    Dequeue(queue, &src, &mesg) ;  
Print_message(src, mesg) ;
```

Termination Detection: The Done function

```
queue_size = enqueued - dequeued ;  
if (queue_size == 0 && done_sending == thread_count)  
    return TRUE;  
else  
    return FALSE;
```



each thread increments this after
completing its *for loop*

Startup (1)

- When the program execution begins, the master thread, will get command line arguments and allocate an array of message queues: one for each thread
- This array will be shared among the threads, since any thread can send to any other thread, and hence any thread can enqueue a message in any of the queues

Startup (2)

- One or more threads may finish allocating their queues before some other threads
- An explicit barrier is needed to blocks all threads in the team that reach the barrier
- Once all threads have reached the barrier, all threads in the team can then proceed

```
#pragma omp barrier
```

Atomic Directive (1)

- Unlike the *critical* directive, it can only protect critical sections that consist of a **single C assignment statement**

pragma omp *atomic*

- Further, the statement must have one of the following forms: **x <op> = <expression>;**

x++;

++x;

x--;

--x;

Atomic Directive (2)

- Here <op> can be one of the binary operators

$+$, $*$, $-$, $/$, $\&$, \wedge , $|$, \ll , or \gg

- Note that <expression> must not reference x

- Example

```
# pragma omp atomic
```

```
 $x += y++$ 
```

a thread's update of x will be completed before any other thread can begin updating x

Atomic Directive (3)

- Many processors provide a special load-modify-store instruction
- A *critical* section that only does a load-modify-store can be protected much more efficiently by using this special instruction rather than the constructs used to protect more general critical sections

Critical (named) Directive

- OpenMP provides the option of adding a **name** to a *critical* directive:

pragma omp *critical* (**name**)

- This allows the simultaneous execution of two blocks protected with *critical* directives with different **names**
- The **names** are set during compilation, so a different *critical* section is needed for each thread's queue

Locks

- A lock consists of a data structure and functions that allow the programmer to explicitly enforce mutual exclusion in a critical section

Locks

```
/* Executed by one thread */  
Initialize the lock data structure ;
```

```
...
```

```
/* Executed by one thread */  
Attempt to lock or set the lock data structure ;  
Critical section ;  
Unlock or unset the lock data structure ;
```

```
...
```

```
/* Executed by one thread */  
Destroy the lock data structure ;
```

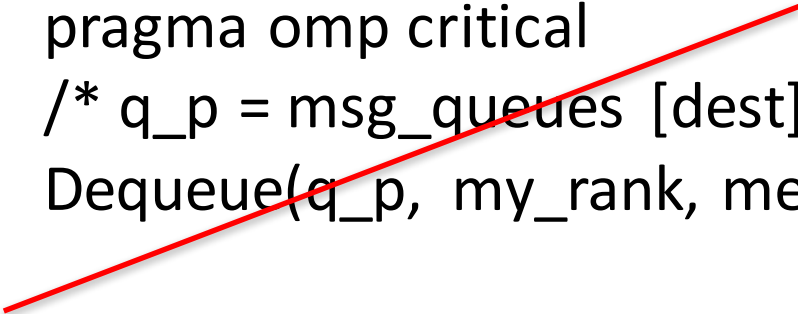
Message Passing Program: Using Locks

```
# pragma omp critical
/* q_p = msg_queues [dest] */
Enqueue(q_p, my_rank, mesg);

/* q_p = msg_queues [dest] */
omp_set_lock(&q_p -> lock) ;
Enqueue(q_p, my_rank, mesg);
omp_unset_lock(&q_p -> lock) ;
```

Message Passing Program: Using Locks

```
# pragma omp critical  
/* q_p = msg_queues [dest] */  
Dequeue(q_p, my_rank, mesg);
```



```
/* q_p = msg_queues [dest] */  
omp_set_lock(&q_p -> lock) ;  
Dequeue(q_p, my_rank, mesg);  
omp_unset_lock(&q_p -> lock) ;
```

Some Caveats

- Do not mix the different types of mutual exclusion for a single critical section
- There is no guarantee of fairness in mutual exclusion constructs
- It can be dangerous to “**nest**” mutual exclusion constructs

Concluding Remarks (1)

- OpenMP is a standard for programming shared-memory systems
- OpenMP uses both special functions and preprocessor directives called pragmas
- OpenMP programs start multiple threads rather than multiple processes
- Many OpenMP directives can be modified by clauses

Concluding Remarks (2)

- A major problem in the development of shared memory programs is the possibility of race conditions
- OpenMP provides several mechanisms for insuring mutual exclusion in critical sections
 - Critical directives
 - Named critical directives
 - Atomic directives
 - Simple locks

Concluding Remarks (3)

- By default most systems use a block-partitioning of the iterations in a parallelized *for loop*
- OpenMP offers a variety of scheduling options
- In OpenMP the scope of a variable is the collection of threads to which the variable is accessible

Concluding Remarks (4)

- A reduction process repeatedly applies the same reduction operator to a sequence of operands in order to get a single result.